

```
In [1]: import pandas as pd
import seaborn as sns
```

```
In [2]: df = pd.read_csv('Churn_Modelling.csv')
```

```
In [3]: df.shape
```

```
Out[3]: (10000, 14)
```

```
In [4]: df.columns
```

```
Out[4]: Index(['RowNumber', 'CustomerId', 'Surname', 'CreditScore', 'Geography',
              'Gender', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard',
              'IsActiveMember', 'EstimatedSalary', 'Exited'],
              dtype='object')
```

```
In [5]: df.head()
```

```
Out[5]:
```


	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	1	15634602	Hargrave	619	France	Female	42	2	0.004563	1	1	1	59330.42	0
1	2	15647311	Hill	608	Spain	Female	41	1	84396.92	1	0	1	26685.55	0
2	3	15619304	Onio	502	France	Female	42	8	89330.63	3	1	1	32695.98	1
3	4	15701354	Boni	699	France	Female	39	1	11352.54	1	0	0	34092.98	0
4	5	15737888	Mitchell	850	Spain	Female	43	2	115966.0	2	0	0	113383.21	0

```
In [6]: # input data
x = df[['CreditScore', 'Age', 'Tenure', 'Balance',
        'NumOfProducts', 'HasCrCard',
        'IsActiveMember', 'EstimatedSalary']]
```

```
# output data
y = df['Exited']
```

```
In [7]: sns.countplot(x = y)
```

```
Out[7]: <Axes: xlabel='Exited', ylabel='count'>
```



```
In [8]: y.value_counts()
```

```
Out[8]: 0    7963
        1    2037
        Name: Exited, dtype: int64
```

```
In [9]: # normalise
        # standardization
        from sklearn.preprocessing import StandardScaler as ss
        scaler = ss()
```

```
In [10]: x_scaled = scaler.fit_transform(x)
```

```
In [11]: x_scaled
```

```
Out[11]: array([[ -0.32622142,  0.29351742, -1.04175968, ...,  0.64609167,
                   0.97024255,  0.02188649],
                 [-0.44003595,  0.19816383, -1.38753759, ..., -1.54776799,
                   0.97024255,  0.21653375],
                 [-1.53679418,  0.29351742,  1.03290776, ...,  0.64609167,
                  -1.03067011,  0.2406869 ],
                 ...,
                 [ 0.60498839, -0.27860412,  0.68712986, ..., -1.54776799,
                   0.97024255, -1.00864308],
                 [ 1.25683526,  0.29351742, -0.69598177, ...,  0.64609167,
                  -1.03067011, -0.12523071],
                 [ 1.46377078, -1.04143285, -0.35020386, ...,  0.64609167,
                  -1.03067011, -1.07636976]])
```

```
In [12]: # cross validation
        from sklearn.model_selection import train_test_split as tts
```

```
In [13]: xtrain,xtest,ytrain,ytest = tts(x_scaled, y, test_size=0.25)
```

```
In [14]: x_scaled.shape
```

```
Out[14]: (10000, 8)
```

```
In [16]: xtest.shape
```

```
Out[16]: (2500, 8)
```

```
In [17]: xtrain.shape
```

```
Out[17]: (7500, 8)
```

```
In [18]: from sklearn.neural_network import MLPClassifier as mlpc
```

```
In [19]: ann = mlpc(hidden_layer_sizes=(100, 100, 100),
                    max_iter=100, activation='relu')
```

```
In [20]: # train
        ann.fit(xtrain, ytrain)
```

C:\Users\sadek\anaconda3\lib\site-packages\sklearn\neural_network_multilayer_perceptron.py:684: ConvergenceWarning: Stochastic Optimizer: Maximum iterations (100) reached and the optimization hasn't converged yet.

warnings.warn(

```
Out[20]: ▼ MLPClassifier
        MLPClassifier(hidden_layer_sizes=(100, 100, 100), max_iter=100)
```

```
In [21]: y_pred = ann.predict(xtest)
```

```
In [22]: y_pred
```

```
Out[22]: array([0, 0, 0, ..., 0, 0, 0], dtype=int64)
```


```
In [23]: from sklearn.metrics import ConfusionMatrixDisplay as matrix
        from sklearn.metrics import classification_report as report
        from sklearn.metrics import accuracy_score as score
```

```
In [24]: ytest.value_counts()
```

```
Out[24]: 0    1983
        1     517
        Name: Exited, dtype: int64
```

```
In [25]: matrix.from_predictions(ytest, y_pred)
```

Out[25]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x2453faf7e20>



In [26]: `score(ytest, y_pred)`

Out[26]: 0.8204

In [27]: `print(report(ytest, y_pred))`

	precision	recall	f1-score	support
0	0.87	0.91	0.89	1983
1	0.58	0.47	0.52	517
accuracy			0.82	2500
macro avg	0.72	0.69	0.70	2500
weighted avg	0.81	0.82	0.81	2500

In [28]: `!pip install imbalanced-learn`

Requirement already satisfied: imbalanced-learn in c:\users\sadek\anaconda3\lib\site-packages (0.10.1)
 Requirement already satisfied: threadpoolctl>=2.0.0 in c:\users\sadek\anaconda3\lib\site-packages (from imbalanced-learn) (2.2.0)
 Requirement already satisfied: joblib>=1.1.1 in c:\users\sadek\anaconda3\lib\site-packages (from imbalanced-learn) (1.1.1)
 Requirement already satisfied: numpy>=1.17.3 in c:\users\sadek\anaconda3\lib\site-packages (from imbalanced-learn) (1.23.5)
 Requirement already satisfied: scikit-learn>=1.0.2 in c:\users\sadek\anaconda3\lib\site-packages (from imbalanced-learn) (1.2.1)
 Requirement already satisfied: scipy>=1.3.2 in c:\users\sadek\anaconda3\lib\site-packages (from imbalanced-learn) (1.10.0)

In [30]: `from imblearn.over_sampling import RandomOverSampler as rs`

In [31]: `ros = rs()`

In [32]: `x_res, y_res = ros.fit_resample(x, y)`

In [33]: `y_res.value_counts()`

Out[33]: `1 7963
0 7963
Name: Exited, dtype: int64`

In [34]: `# normalize
from sklearn.preprocessing import StandardScaler as ss`

In [35]: `scaler = ss()`

In [38]: `x_scaled = scaler.fit_transform(x_res)`

In [39]: `x_scaled`

Out[39]: `array([[-0.29941334, 0.07851204, -1.02461935, ..., 0.64646199,
 1.08538632, 0.0183093],
 [-0.41136471, -0.01571424, -1.37009269, ..., -1.54688135,
 1.08538632, 0.21233705],
 [-1.49016885, 0.07851204, 1.0482207 , ..., 0.64646199,
 -0.92133094, 0.23641332],
 ...,
 [1.2068415 , 0.36119087, 0.70274735, ..., 0.64646199,
 -0.92133094, -1.27686147],
 [0.27052093, -0.76952445, -1.71556604, ..., 0.64646199,
 -0.92133094, -0.61495554],
 [1.18648671, 0.17273831, -0.33367267, ..., -1.54688135,
 1.08538632, 1.2755192]])`

In [40]: `x_scaled.shape`

Out[40]: `(15926, 8)`

In [41]: `# cross validation`

```
In [66]: xtrain,xtest,ytrain,ytest = tts(x_scaled, y_res, test_size=0.3)
```

```
In [67]: x_res.shape
```

```
Out[67]: (15926, 8)
```

```
In [68]: xtest.shape
```

```
Out[68]: (3186, 8)
```

```
In [69]: xtrain.shape
```

```
Out[69]: (12740, 8)
```

```
In [70]: ann.fit(xtrain, ytrain)
```

```
C:\Users\sadek\anaconda3\lib\site-packages\sklearn\neural_network\_multilayer_perceptron.py:684: ConvergenceWarning: Stochastic Optimizer: Maximum iterations (100) reached and the optimization hasn't converged yet.
  warnings.warn(
```

```
Out[70]: MLPClassifier
MLPClassifier(hidden_layer_sizes=(100, 100, 100), max_iter=100)
```


```
In [71]: y_pred = ann.predict(xtest)
```

```
In [72]: ytest.value_counts()
```

```
Out[72]: 0    1603
         1    1583
         Name: Exited, dtype: int64
```

```
In [73]: matrix.from_predictions(ytest, y_pred)
```

```
Out[73]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x245454b1ea0>
```



```
In [74]: print(score(ytest, y_pred))
```

```
0.8747645951035782
```

```
In [75]: print(report(ytest, y_pred))
```

	precision	recall	f1-score	support
0	0.91	0.83	0.87	1603
1	0.84	0.92	0.88	1583
accuracy			0.87	3186
macro avg	0.88	0.88	0.87	3186
weighted avg	0.88	0.87	0.87	3186

```
In [76]: from imblearn.under_sampling import RandomUnderSampler as rs
```

```
In [77]: rus = rs()
```

```
In [78]: x_res, y_res = rus.fit_resample(x, y)
```

```
In [79]: x_scaled = scaler.fit_transform(x_res)
```

```
In [80]: xtrain,xtest,ytrain,yest = tts(x_scaled, y_res, test_size=0.2)
```

```
In [86]: ann.fit(xtrain,ytrain)
```

```
C:\Users\sadek\anaconda3\lib\site-packages\sklearn\neural_network\_multilayer_perceptron.py:684: ConvergenceWarning: Stochastic Optimizer: Maximum iterations (100) reached and the optimization hasn't converged yet.  
  warnings.warn(
```

Out[86]:

```
▼ MLPClassifier  
MLPClassifier(hidden_layer_sizes=(100, 100, 100), max_iter=100)
```

In [87]: `y_pred = ann.predict(xtest)`In [88]: `matrix.from_predictions(ytest, y_pred)`


```

-----
ValueError                                Traceback (most recent call last)
Cell In[88], line 1
----> 1 matrix.from_predictions(ytest, y_pred)

File ~\anaconda3\lib\site-packages\sklearn\metrics\_plot\confusion_matrix.py:463, in ConfusionMatrixDisplay.from_predictions(cls, y_true, y_pred, labels, sample_weight, normalize, display_labels, include_values, xticks_rotation, values_format, cmap, ax, colorbar, im_kw, text_kw)
    460     else:
    461         display_labels = labels
--> 463 cm = confusion_matrix(
    464     y_true,
    465     y_pred,
    466     sample_weight=sample_weight,
    467     labels=labels,
    468     normalize=normalize,
    469 )
    471 disp = cls(confusion_matrix=cm, display_labels=display_labels)
    473 return disp.plot(
    474     include_values=include_values,
    475     cmap=cmap,
    (...)
    481     text_kw=text_kw,
    482 )

File ~\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:317, in confusion_matrix(y_true, y_pred, labels, sample_weight, normalize)
    232 def confusion_matrix(
    233     y_true, y_pred, *, labels=None, sample_weight=None, normalize=None
    234 ):
    235     """Compute confusion matrix to evaluate the accuracy of a classification.

    236
    237     By definition a confusion matrix :math:`C` is such that :math:`C_{ij}`
    (...)
    315     (0, 2, 1, 1)
    316     """
--> 317     y_type, y_true, y_pred = _check_targets(y_true, y_pred)
    318     if y_type not in ("binary", "multiclass"):
    319         raise ValueError("%s is not supported" % y_type)

File ~\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:86, in _check_targets(y_true, y_pred)
    59 def _check_targets(y_true, y_pred):
    60     """Check that y_true and y_pred belong to the same classification task.
    61
    62     This converts multiclass or binary types to a common shape, and raises a
    (...)
    84     y_pred : array or indicator matrix
    85     """
--> 86     check_consistent_length(y_true, y_pred)
    87     type_true = type_of_target(y_true, input_name="y_true")
    88     type_pred = type_of_target(y_pred, input_name="y_pred")

File ~\anaconda3\lib\site-packages\sklearn\utils\validation.py:397, in check_consist

```

```
ent_length(*arrays)
    395 uniques = np.unique(lengths)
    396 if len(uniques) > 1:
--> 397     raise ValueError(
    398         "Found input variables with inconsistent numbers of samples: %r"
    399         % [int(1) for l in lengths]
    400     )
```

ValueError: Found input variables with inconsistent numbers of samples: [3186, 815]

In [84]: `y_res.value_counts()`

Out[84]:

0	2037
1	2037

Name: Exited, dtype: int64

In []: